

HW6 成果分析

```
//TODO
```

```
Fix bugs of Batchnorm backprop
```

本次作業實作CNN神經網路的建構

架構

Convolutional Layer

先建立最主要的卷積層。根據輸入輸出的 dimensions 還有 kernal 的設定，我們可以用巢狀的迴圈進行卷積。

```
# Forward

for nn in range(x_shape[0]):
    for ff in range(w_shape[0]):
        for hh in range(output_H):
            for ww in range(output_W):
                conv_window_h = hh * stride
                conv_window_w = ww * stride
                out_tmp = x_pad[nn, :, conv_window_h:conv_window_h+w_shape[2],
conv_window_w:conv_window_w+w_shape[3]] * w[ff, :, :, :]
                out[nn, ff, hh, ww] = torch.sum(out_tmp) + b[ff]
```

我的作法是先框出一個 kernal window，後續的所有運算只要專注在這個 window 當中即可。

```
#Backward

for nn in range(x_shape[0]):
    for ff in range(w_shape[0]):
        for hh in range(output_H):
            for ww in range(output_W):
                conv_window_h = hh * stride
                conv_window_w = ww * stride
                dx_pad[nn, :, conv_window_h:conv_window_h+w_shape[2],
conv_window_w:conv_window_w+w_shape[3]] += dout[nn, ff, hh, ww] * w[ff, :, :, :]
                dw[ff, :, :, :] += dout[nn, ff, hh, ww] * x_pad[nn, :,
conv_window_h:conv_window_h+w_shape[2],
conv_window_w:conv_window_w+w_shape[3]]
                db[ff] += dout[nn, ff, hh, ww]

dx = dx_pad[:, :, pad:pad+x_shape[2], pad:pad+x_shape[3]]
```

在反向傳播中，因為 $(out=x*w)$ 所以 $(dx=dout*dw)$ 。又因為有 rolling windows 的關係，所以矩陣內的每個元素都有可能被重複運算到，故使用 `+=` 運算子，將所有的變量疊加起來。

Maxpooling

和上面同理，先專注在一個 pooling window 的運算上，再用巢狀迴圈 rolling。Maxpooling 是在一個 pooling window 中，輸出一個最大值。

```
# Forward
for nn in range(x_shape[0]):
    for cc in range(x_shape[1]):
        for hh in range(output_H):
            for ww in range(output_W):
                pool_window_h = hh * stride
                pool_window_w = ww * stride
                out_tmp = x[nn, cc, pool_window_h:pool_window_h+pool_height,
pool_window_w:pool_window_w+pool_width]
                out[nn, cc, hh, ww] = torch.max(out_tmp)

# Backward

for nn in range(x_shape[0]):
    for cc in range(x_shape[1]):
        for hh in range(output_H):
            for ww in range(output_W):
                pool_window_h = hh * stride
                pool_window_w = ww * stride
                out_tmp = x[nn, cc, pool_window_h:pool_window_h+pool_height,
pool_window_w:pool_window_w+pool_width]
                out_tmp_max = torch.max(out_tmp)
                max_mat = (out_tmp == out_tmp_max)
                dx[nn, cc, pool_window_h:pool_window_h+pool_height,
pool_window_w:pool_window_w+pool_width] += (dout[nn, cc, hh, ww]*max_mat)
```

Maxpooling window 中，梯度只會從最大值的元素流過，因為輸出只和輸入的最大值有關係。

$$\begin{cases} \text{if } \max(\mathbf{x}) = x_i & \text{then } dx_i = dout \\ \text{otherwise} & dx_i = 0 \end{cases}$$

在反向傳播中，又運算了一次 max，其實可以在正向傳播完成時直接把 out 存到 cache，增加效率。(不過比大小應該耗不掉多少效率)

Batch Normalization

以 mini-batch 為單位，依照各個 mini-batch 來進行正規化。如此便可以減少模型過度依賴預設值。

- 調整最佳化過程可能只對最後幾層有影響，因為神經網路的複雜度讓梯度傳至前面可能已經使特徵發生變化。
- 確保每一層的特徵相同，所以在每一層結束後都進行正規化。

根據論文：

$$\hat{x}^{(k)} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

反向傳播：

$$\begin{aligned}\frac{\partial \ell}{\partial \hat{x}_i} &= \frac{\partial \ell}{\partial y_i} \cdot \gamma \\ \frac{\partial \ell}{\partial \sigma_B^2} &= \sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot (x_i - \mu_B) \cdot \frac{-1}{2} (\sigma_B^2 + \epsilon)^{-3/2} \\ \frac{\partial \ell}{\partial \mu_B} &= \left(\sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma_B^2 + \epsilon}} \right) + \frac{\partial \ell}{\partial \sigma_B^2} \cdot \frac{\sum_{i=1}^m -2(x_i - \mu_B)}{m} \\ \frac{\partial \ell}{\partial x_i} &= \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma_B^2 + \epsilon}} + \frac{\partial \ell}{\partial \sigma_B^2} \cdot \frac{2(x_i - \mu_B)}{m} + \frac{\partial \ell}{\partial \mu_B} \cdot \frac{1}{m} \\ \frac{\partial \ell}{\partial \gamma} &= \sum_{i=1}^m \frac{\partial \ell}{\partial y_i} \cdot \hat{x}_i \\ \frac{\partial \ell}{\partial \beta} &= \sum_{i=1}^m \frac{\partial \ell}{\partial y_i}\end{aligned}$$

加入我們的 DNN 當中。

```
Before batch normalization:
means: ['52.046', '11.122', '10.243']
stds:  ['34.646', '30.732', '39.429']

After batch normalization (gamma=1, beta=0)
means: ['-0.000', '-0.000', '-0.000']
stds:  ['1.000', '1.000', '1.000']

After batch normalization (gamma= [1.0, 2.0, 3.0] , beta= [11.0, 12.0, 13.0] )
means: ['11.000', '12.000', '13.000']
stds:  ['1.000', '2.000', '3.000']
```

可以發現，在加入 Batchnorm 之後，每層的標準差和平均數都可以被很好的控制。

```
Running check with reg = 0
W1 max relative error: 7.554828e-03
W2 max relative error: 9.331857e-03
W3 max relative error: 1.078155e-02
W4 max relative error: 1.501991e-09
b1 max relative error: 9.999999e-01
b2 max relative error: 1.000000e+00
b3 max relative error: 1.110223e-06
b4 max relative error: 2.382429e-09
beta1 max relative error: 7.290488e-03
beta2 max relative error: 2.329975e-02
beta3 max relative error: 2.572281e-09
gamma1 max relative error: 6.621865e-03
gamma2 max relative error: 2.271405e-02
gamma3 max relative error: 1.686124e-09

Running check with reg = 3.14
W1 max relative error: 1.145399e-03
W2 max relative error: 1.503421e-03
W3 max relative error: 1.830862e-03
W4 max relative error: 2.206159e-08
b1 max relative error: 1.942890e-06
b2 max relative error: 1.804112e-06
b3 max relative error: 1.110223e-06
b4 max relative error: 9.149199e-08
beta1 max relative error: 6.879257e-03
beta2 max relative error: 7.083755e-03
beta3 max relative error: 6.922886e-08
gamma1 max relative error: 6.652254e-03
gamma2 max relative error: 8.533660e-03
gamma3 max relative error: 4.155353e-08
```

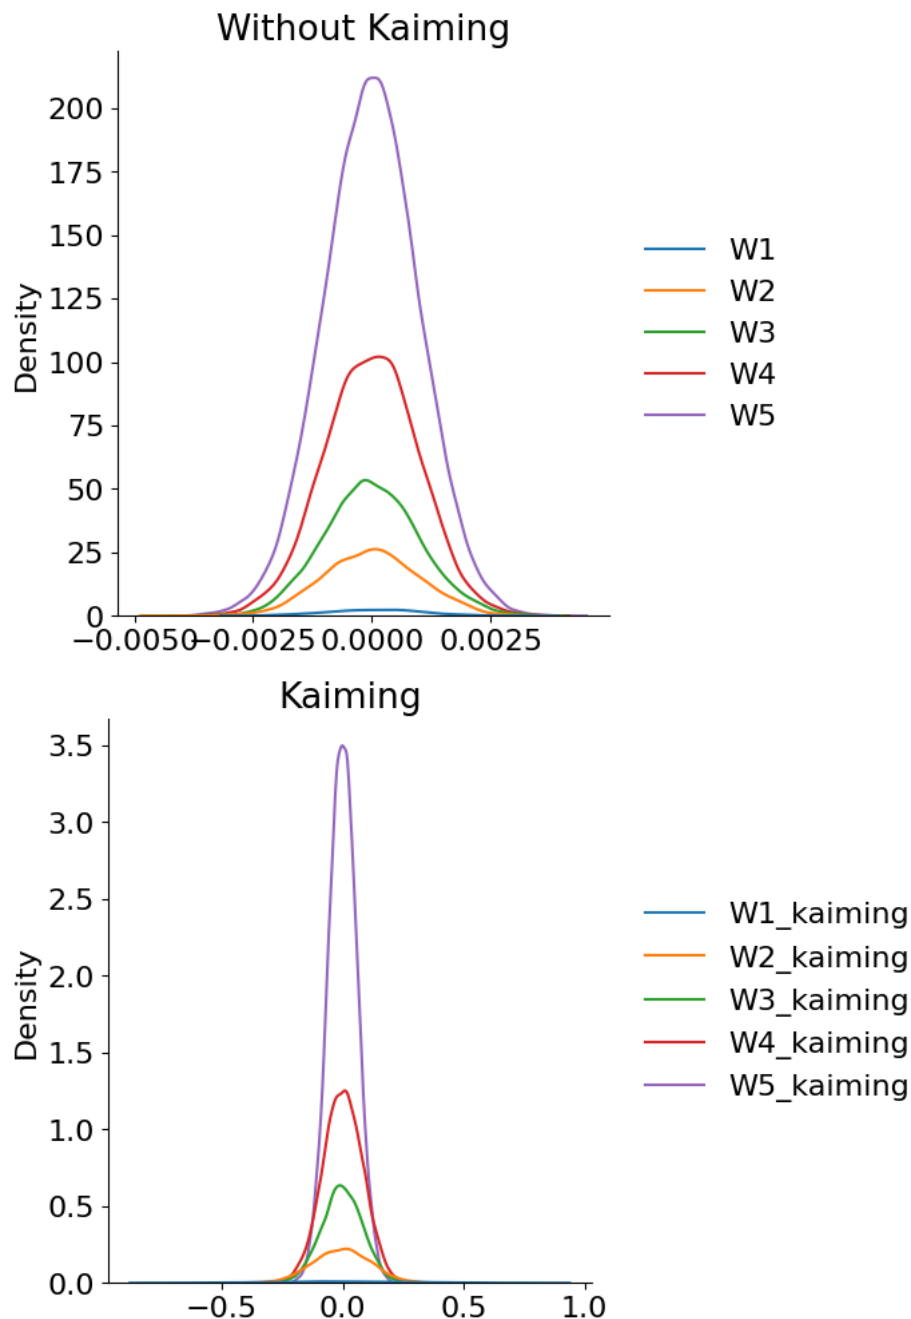
只是不知道為何，算出來的結果一直有誤差，檢查了幾次都發現不出問題，上網查閱相關的範例也是一樣。日後多查一些資料後再進行改善。

backprop of BN:

```
dx error: 0.07613132307158305
dgamma error: 3.087149574058185e-10
dbeta error: 2.919336030277982e-10
```

Kaiming Initialization

- ReLU 活化函數會拋棄小於0的值
- 若每層的權重是常態分佈 ($\mu=0$) 的話，等於說砍掉了一半的值。
- 我的理解是，Kaiming 針對這個分佈曲線做了平移和縮放，讓權重有不同的變化



3 Layers CNN

實作三層的 cnn 神經網路，採用剛才的 API，並且新增 ReLU 的運算就好了。

```
# ReLU  
  
out = torch.max(torch.zeros_like(out), out) # Forward  
dout = (dout >= torch.zeros_like(dout)).long() # Backward
```

```

# Forward

out, cache1 = Conv.forward(X, W1, b1, conv_param)
cache2 = out
out = torch.max(torch.zeros_like(out), out)
out, cache3 = MaxPool.forward(out, pool_param)
out, cache4 = Linear_ReLU.forward(out, W2, b2)
scores, cache5 = Linear.forward(out, W3, b3)

# Backward

loss, dout = softmax_loss(scores, y)
loss += self.reg * ((torch.sum(W1 * W1) + torch.sum(W2 * W2) +
torch.sum(W3 * W3)))

dh, grads['W3'], grads['b3'] = Linear.backward(dout, cache5)
grads['W3'] += 2 * self.reg * W3
dh, grads['W2'], grads['b2'] = Linear_ReLU.backward(dh, cache4)
grads['W2'] += 2 * self.reg * W2
dh = MaxPool.backward(dh, cache3)
dh[cache2 < 0] = 0
_, grads['W1'], grads['b1'] = Conv.backward(dh, cache1)
grads['W1'] += 2 * self.reg * W1

```

事實上，也可以改用 `torch.nn.conv2d` 平行處理，大幅增加運算速度。直接引用下方的 FastConv API 就可以了。

Deeper CNN

用 for 處理深層神經網路。

```

# Forward

for i in range(self.num_layers-1):

    # Conv forward
    h, cache = FastConv.forward(h, self.params['W'+str(i+1)],
self.params['b'+str(i+1)], conv_param)
    caches.append(cache)

    # BN forward
    if self.batchnorm:
        h, cache = SpatialBatchNorm.forward(h, self.params['gamma'+str(i+1)],
self.params['beta'+str(i+1)], self.bn_params[i])
        caches.append(cache)

    # ReLU forward
    h, cache = ReLU.forward(h)
    caches.append(cache)

    # Pool forward
    if i in self.max_pools:
        h, cache = FastMaxPool.forward(h, pool_param)
        caches.append(cache)

    scores, cache_final = Linear.forward(h,
self.params['W'+str(self.num_layers)], self.params['b'+str(self.num_layers)])

```

```

# Backward

loss, dout = softmax_loss(scores, y)
loss += self.reg * (torch.sum(self.params['W'+str(self.num_layers)]**2))

dx, dw, db = Linear.backward(dout, cache_final)
grads['W'+str(self.num_layers)] = dw + (2 * self.reg *
self.params['W'+str(self.num_layers)])
grads['b'+str(self.num_layers)] = db

for i in range(self.num_layers-1, 0, -1):
    #print(str(i)+"---")

    # Pool backward
    if i-1 in self.max_pools:
        dx = FastMaxPool.backward(dx, caches[-1])
        #print(len(caches[-1]))
        caches.pop()

    # ReLU backward
    dx = ReLU.backward(dx, caches[-1])
    caches.pop()

    # BN backward
    if self.batchnorm:
        dx, dgamma, dbeta = SpatialBatchNorm.backward(dx, caches[-1])
        caches.pop()
        grads['gamma'+str(i)] = dgamma
        grads['beta'+str(i)] = dbeta

    # Conv backward
    dx, dw, db = FastConv.backward(dx, caches[-1])

    #print(len(caches[-1]))
    caches.pop()
    grads['W'+str(i)] = dw + (2 * self.reg * self.params['W'+str(i)])
    grads['b'+str(i)] = db
    loss += self.reg * (torch.sum(self.params['W'+str(i)]**2))

```

這部份做了很多次的調整，才形成現在這個版本。
 最一開始在做反向傳播時是直接取用 `cache[i]`，但有些層不需經過池化，所以造成在正向傳播時沒有池化的層我也必須 `append` 一個佔位用的假 `cache` 給他，非常不漂亮，到後來增加了 `Batchnorm` 之後更是如此。
 所以，最後才改成以上的版本，反向傳播中用完一個 `cache` 就 `pop` 掉，每次讀取 `cache[-1]` 即可。

模型訓練

Train the best convolutional model that you can on CIFAR-10, storing your best model in the `best_model` variable. We require you to get at least 71% accuracy on the validation set using a convolutional net, within 60 seconds of training.

在一分鐘訓練出至少71%準確度的模型。以下是我的嘗試：

第一個版本：

```
#from fc_networks import adam, sgd_momentum

weight_scale = 'kaiming'
learning_rate = 0.0035
reg = 0.001

model = DeepConvNet(input_dims=(3, 32, 32), num_classes=10,
                    num_filters=[32, 80],
                    max_pools=[0, 1],
                    reg=reg,
                    weight_scale=weight_scale,
                    dtype=dtype,
                    device=device)
solver = Solver(model, data_dict,
               num_epochs=30,
               batch_size=128,
               update_rule=adam,
               optim_config={
                   'learning_rate': learning_rate,
               },
               print_every=10000,
               device=device)
```

```
val_acc = solver.check_accuracy(data_dict['X_val'], data_dict['y_val'])
test_acc = solver.check_accuracy(data_dict['X_test'], data_dict['y_test'])

print(f'Validation set accuracy: {("{:.4f}".format(val_acc*100))}%')
print(f'Test set accuracy: {("{:.4f}".format(test_acc*100))}%')
```

Validation set accuracy: 71.1700%
Test set accuracy: 70.3900%

第二個版本：

增加神經數量， capacity 變高，所以效果變好，但如果仔細看訓練過程，其實是 overfitting 的。

```

weight_scale = 'kaiming'
learning_rate = 0.0025
reg = 0.001

model = DeepConvNet(input_dims=(3, 32, 32), num_classes=10,
                    num_filters=[32, 128],
                    max_pools=[0, 1],
                    reg=reg,
                    weight_scale=weight_scale,
                    # batchnorm=True,
                    dtype=dtype,
                    device=device)
solver = Solver(model, data_dict,
                num_epochs=30,
                batch_size=128,
                update_rule=adam,
                optim_config={
                    'learning_rate': learning_rate,
                },
                print_every=10000,
                device=device)

```

```

▶ 1 val_acc = solver.check_accuracy(data_dict['X_val'], data_dict['y_val'])
2 test_acc = solver.check_accuracy(data_dict['X_test'], data_dict['y_test'])
3
4 print(f'Validation set accuracy: {"{:.4f}".format(val_acc*100)}%')
5 print(f'Test set accuracy: {"{:.4f}".format(test_acc*100)}%')

```

Validation set accuracy: 71.5700%
Test set accuracy: 71.5200%

```

(Time 0.01 sec; Iteration 1 / 9360) loss: 2.712113
(Epoch 0 / 30) train acc: 10.90%; val_acc: 12.28%
(Epoch 1 / 30) train acc: 60.80%; val_acc: 59.32%
(Epoch 2 / 30) train acc: 66.70%; val_acc: 63.31%
(Epoch 3 / 30) train acc: 70.20%; val_acc: 66.83%
(Epoch 4 / 30) train acc: 73.20%; val_acc: 68.54%
(Epoch 5 / 30) train acc: 72.00%; val_acc: 68.26%
(Epoch 6 / 30) train acc: 71.80%; val_acc: 69.91%
(Epoch 7 / 30) train acc: 74.80%; val_acc: 69.54%
(Epoch 8 / 30) train acc: 74.40%; val_acc: 69.44%
(Epoch 9 / 30) train acc: 76.90%; val_acc: 70.68%
(Epoch 10 / 30) train acc: 74.60%; val_acc: 69.40%
(Epoch 11 / 30) train acc: 75.80%; val_acc: 70.76%
(Epoch 12 / 30) train acc: 75.90%; val_acc: 69.35%
(Epoch 13 / 30) train acc: 77.40%; val_acc: 70.00%
(Epoch 14 / 30) train acc: 77.50%; val_acc: 68.26%
(Epoch 15 / 30) train acc: 77.10%; val_acc: 70.59%
(Epoch 16 / 30) train acc: 75.00%; val_acc: 69.11%
(Epoch 17 / 30) train acc: 79.70%; val_acc: 71.57%
(Epoch 18 / 30) train acc: 75.90%; val_acc: 70.72%
(Epoch 19 / 30) train acc: 79.90%; val_acc: 71.43%
(Epoch 20 / 30) train acc: 79.80%; val_acc: 71.28%
(Epoch 21 / 30) train acc: 78.00%; val_acc: 71.07%
(Epoch 22 / 30) train acc: 80.20%; val_acc: 71.10%
(Time 60.00 sec; Iteration 7158 / 9360) loss: 0.991407
End of training; next iteration will exceed the time limit.

```


第三個版本：

加入 Batchnorm

```
weight_scale = 'kaiming'
learning_rate = 0.0025
reg = 0.001

model = DeepConvNet(input_dims=(3, 32, 32), num_classes=10,
                    num_filters=[32, 100],
                    max_pools=[0, 1],
                    reg=reg,
                    weight_scale=weight_scale,
                    batchnorm=True,
                    dtype=dtype,
                    device=device)
solver = Solver(model, data_dict,
                num_epochs=30,
                batch_size=128,
                update_rule=adam,
                optim_config={
                    'learning_rate': learning_rate,
                },
                print_every=10000,
                device=device)
```


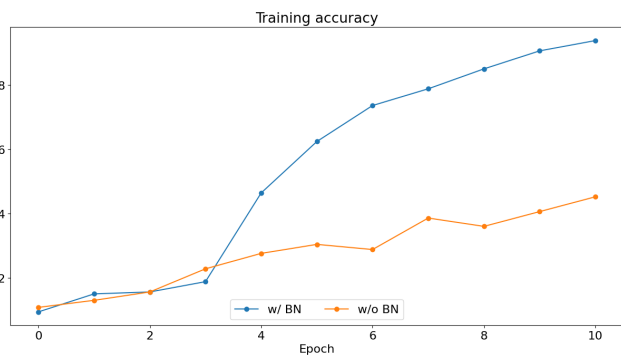
```
(Time 0.02 sec; Iteration 1 / 9360) loss: 3.577355
(Epoch 0 / 30) train acc: 10.70%; val_acc: 11.50%
(Epoch 1 / 30) train acc: 61.40%; val_acc: 58.19%
(Epoch 2 / 30) train acc: 68.40%; val_acc: 64.61%
(Epoch 3 / 30) train acc: 71.20%; val_acc: 66.40%
(Epoch 4 / 30) train acc: 71.80%; val_acc: 66.55%
(Epoch 5 / 30) train acc: 71.60%; val_acc: 66.70%
(Epoch 6 / 30) train acc: 76.00%; val_acc: 68.18%
(Epoch 7 / 30) train acc: 79.80%; val_acc: 70.36%
(Epoch 8 / 30) train acc: 79.00%; val_acc: 68.78%
(Epoch 9 / 30) train acc: 74.70%; val_acc: 67.19%
(Epoch 10 / 30) train acc: 80.10%; val_acc: 69.83%
(Epoch 11 / 30) train acc: 82.00%; val_acc: 70.29%
(Epoch 12 / 30) train acc: 79.60%; val_acc: 70.23%
(Time 59.99 sec; Iteration 3985 / 9360) loss: 0.696833
End of training; next iteration will exceed the time limit.
```

同樣的 `reg` 下，過度擬和的現象減少，但可能是演算法錯誤的關係，讓 testing data 中的準確率極低。

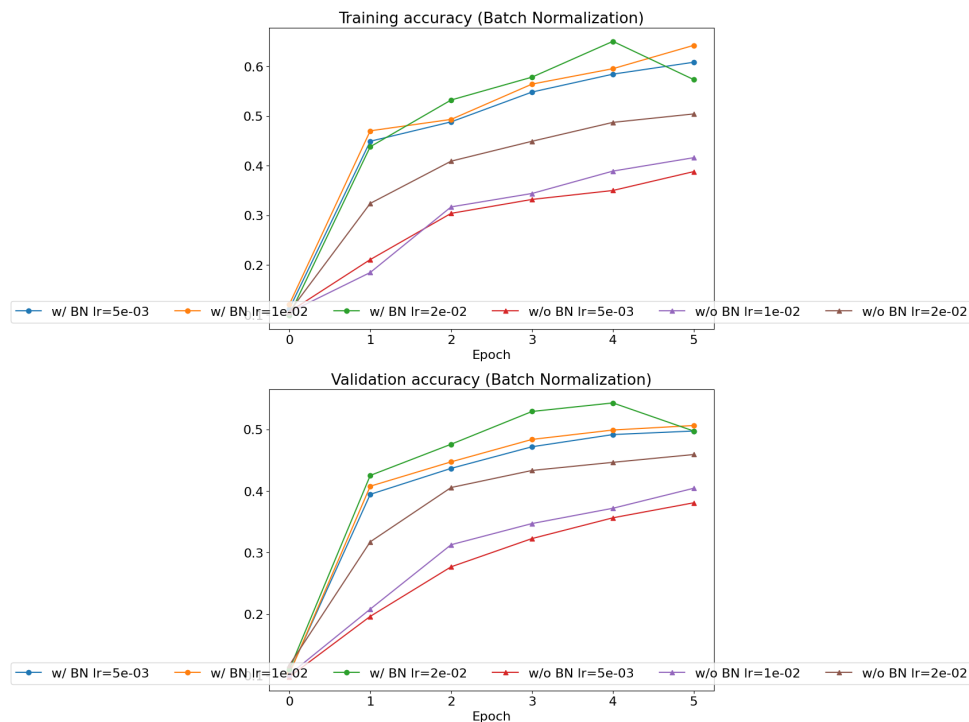
```
Validation set accuracy: 37.7600%
Test set accuracy: 38.3500%
```

我最後採用第二個版本。

訓練結果分析



Batchnorm 在訓練上有很大的提升，但如同上方所述，可能在算法中出現一些錯誤導致 testing 準確率極低。



發現最好的組合是 lr=2e-2 w/BN (表中綠色曲線)，只是到最

後 overfitting，可以用 Early stop 避免準確率下降。
若要穩定成長的話，可以稍微降低學習率至 $lr=1e-2$ (表中橘色曲線)。

其他分析

為什麼 Regularization 可以防止過度擬合

$$\text{Cost Function } C = L(\theta) + \lambda\Omega(W)$$

$$\frac{\partial C}{\partial W} = \frac{\partial L(\theta)}{\partial W} + \lambda \frac{\partial \Omega(W)}{\partial W}$$

$$\frac{\partial C}{\partial b} = \frac{\partial L(\theta)}{\partial b} + 0$$

L2 正則化只會對 W 產生影響。

$$\lambda \frac{\partial \Omega(W)}{\partial w_i} = \frac{\lambda}{n} w_i$$

$$w_i \leftarrow w_i - \eta \frac{\partial C}{\partial w_i}$$

$$= w_i - \eta \frac{\partial L(\theta)}{\partial w_i} - \eta \frac{\lambda}{n} w_i$$

$$= (1 - \eta \frac{\lambda}{n}) w_i - \eta \frac{\partial L(\theta)}{\partial w_i}$$

其中， $(1 - \eta \frac{\lambda}{n}) < 1$ ，所以會讓權重減小，不會放大而讓模型發散。這就是為何正則化可以防止過度擬合。

我們畫個圖證明一下：

